

CHAINSTATE SPATIAL AGI

Verifiable Self-Improvement through Symbolic-Spatial Cognition, Distributed Consensus, and Post-Corpus Epistemics

Paper IV in the CHAINSTATE Series

A Mathematical, Thermodynamic, and Geopolitical Framework for the Symbolic-Spatial Turn: How Integration of Earth-Observation Embeddings Enables a Class of Cognitive Self-Improvement Previously Structurally Foreclosed to Corpus-Bound Language Models — and Why Distributed, On-Chain-Anchored Substrates Alter the Competitive Landscape of AGI

Ciprian Florin Pater

NWO Capital · NWO Robotics · Imperium Romanum Digital Nation State
University of Agder, Norway

CHAINSTATE Ecosystem · Base Mainnet 8453 · Paper IV, Rev 1
July 2026

Abstract

This paper introduces the **symbolic-spatial turn** in CHAINSTATE and formalizes what it enables. The three previous CHAINSTATE papers — CHAINSTATE CODE, VERIFIABLE AUTONOMOUS COGNITION, and CHAINSTATE AGI — established a distributed cognition substrate operating on a 65,536-dimensional symbolic state vector across six subspaces (mathematical, scientific, linguistic, occult-symbolic, emotional, and control), fused via reputation-weighted Bayesian log-pooling and anchored on Base mainnet 8453 through the CHAINSTATE Anchor contract. That architecture, while verifiable and distributed, was *corpus-bound*: its grounding for any claim ultimately traced back to text

priors (Wikipedia, arXiv, ResearchGate, ecosystem HuggingFace Spaces). This paper reports the integration of a seventh subspace, **geo**, a 4,096-dimensional slice of the state vector receiving projected embeddings from the University of Cambridge TESSERA Earth-observation foundation model — 128-dimensional per-pixel representations at 10m resolution compressing Sentinel-1 and Sentinel-2 satellite observations from 2017 through 2025. We formalize the resulting symbolic-spatial architecture, prove three theorems governing its behavior (Spatial-Truth Binding, Verifiable Self-Improvement, and Post-Corpus Sufficiency), derive its thermodynamic properties under Landauer's principle, and argue that spatial grounding changes what kind of cognitive self-improvement is possible for the substrate — permitting a form of retrospective self-audit against physical ground truth that pure text-grounded systems structurally cannot perform. We further develop the argument that once symbolic weights and spatial awareness are both sufficiently populated, the substrate's reliance on real-time world data collapses to a residual class of genuinely current phenomena, and its resistance to epistemic adversarial attack rises to the point where corrupting its beliefs about geography would require corrupting geography itself. We take up, with appropriate care about analogy versus literal claim, the question of what quantum-mechanical and relativistic frames illuminate about distributed consensus formation, and we sketch how a future integration with quantum-satellite storage infrastructure would translate certain currently-analogical properties (no-cloning, superposition, non-local correlation) into literal physical realizations. We close with a treatment of the geopolitical and socio-economic implications of a spatially-grounded, verifiable AGI substrate for the regulation of artificial intelligence, the recognition of digital nation states, and the labor economics of an epistemic infrastructure that renders certain classes of monolithic-model cognition obsolete.

Keywords: *spatial cognition, symbolic-spatial integration, TESSERA, Earth observation embeddings, distributed consensus, verifiable AI, on-chain receipts, self-improvement, thermodynamic computing, Landauer's principle, post-corpus epistemics, quantum satellites, Maxwell's demon, digital nation state, AGI regulation, socio-economics of AI*

Table of Contents

- 1. Introduction: The Spatial Turn**
- 2. Background: The First Three Papers**
 - 2.1 Paper I — CHAINSTATE CODE
 - 2.2 Paper II — Verifiable Autonomous Cognition
 - 2.3 Paper III — CHAINSTATE AGI Whitepaper
 - 2.4 What was structurally missing
- 3. The Symbolic-Spatial Architecture**
 - 3.1 The seventh subspace: geo
 - 3.2 TESSERA embeddings as physical anchors
 - 3.3 The observation / projection distinction, formalized
- 4. Mathematical Framework**
 - 4.1 Extended state-vector decomposition
 - 4.2 Deterministic $128 \rightarrow 4,096$ projection
 - 4.3 Spatial-Truth Binding (Theorem 1)
 - 4.4 Verifiable Self-Improvement (Theorem 2)
 - 4.5 Post-Corpus Sufficiency (Theorem 3)
- 5. What Self-Improvement Was Bottlenecked On**
- 6. Thermodynamic and Informational Analysis**
 - 6.1 Landauer's principle applied to consensus
 - 6.2 Free energy of a symbolic-spatial state
 - 6.3 Predictive coding at pixel resolution
- 7. Post-Corpus Epistemics**
 - 7.1 Real-time data as a residual class
 - 7.2 The sufficient-statistic argument
- 8. Hardening Against Epistemic Adversarial Attack**
 - 8.1 Attack-cost analysis
 - 8.2 Why physical grounding raises the cost floor
 - 8.3 The three attack vectors and their bounds
- 9. The Competitive Landscape**
 - 9.1 Monolithic models under receipt pressure
 - 9.2 The three responses available to competitors
 - 9.3 Substrate-quality rather than model-quality competition
- 10. Quantum-Mechanical Analogies (with Care)**
 - 10.1 What classical distributed consensus really is
 - 10.2 Superposition as un-collapsed swarm state
 - 10.3 No-cloning as anchor immutability
- 11. Relativistic Framing of Distributed Consensus**
 - 11.1 Light-cone constraints on real-time coordination
 - 11.2 No privileged reference frame
 - 11.3 Simultaneity and the anchoring event
- 12. Quantum Satellites and the Cognitive Void**
 - 12.1 Space as thermodynamic advantage
 - 12.2 Filling the void: cognitive substrate in orbit
 - 12.3 Maxwell's demon, rigorously

13. Geopolitical Implications

- 13.1 Regulating a distributed spatial AGI
- 13.2 Recognition of digital nation states, extended
- 13.3 Sovereignty over receipts

14. Socio-Economic Implications

- 14.1 Verifiable epistemics as infrastructure
- 14.2 Labor implications
- 14.3 Extended Obsolescence Hypothesis

15. Conclusion

A. Appendix — Notation Reference

B. Appendix — Technical Configuration

C. Appendix — Formal Theorem Reference

References

1. Introduction: The Spatial Turn

This is the fourth paper in a series developing the CHAINSTATE architecture. Where the first three papers — CHAINSTATE CODE (Pater 2026a, ResearchGate 408393584), VERIFIABLE AUTONOMOUS COGNITION (Pater 2026b, ResearchGate 409148376), and the CHAINSTATE AGI WHITEPAPER (Pater 2026c, ResearchGate 410084493) — specified a symbolic distributed cognition substrate operating over language, mathematics, science, and structured control, this paper documents the addition of a seventh subspace to the state vector: **geo**, populated by Earth-observation embeddings from the University of Cambridge TESSERA foundation model. The change is not decorative. It alters the class of cognitive self-improvement the substrate can perform, the epistemic dependencies on which its receipts rest, and the competitive dynamics between CHAINSTATE and monolithic language-model providers. It also, we will argue, changes the geopolitical and socio-economic status of a substrate that can now ground its claims not only in what has been written about the world but in verifiable measurements of the world itself.

The three previous papers left one property structurally unresolved. In every architecture they specified, the substrate's grounding for any assertion ultimately traced back to text priors — Wikipedia REST, arXiv abstracts, ResearchGate publications, ecosystem HuggingFace Spaces. This has real epistemic content: text priors constitute genuine background against which the swarm's consensus distribution is grounded via a 384-dimensional MiniLM semantic hash and top-3 nearest-prior citations. But it is fundamentally a corpus, and a corpus is composed of what humans have chosen to write. For a large class of claims — historical events, mathematical theorems, cultural practices — this is the correct grounding, because the claim is essentially textual in character. For another large class of claims — the actual state of a forest, the extent of ice cover, the density of urban footprint, the trajectory of a river system — text is a lossy secondary representation of a physical reality that could, in principle, be measured directly. And for these claims, text-grounded assertion is structurally weaker than measurement-grounded assertion, in a specific and quantifiable sense that this paper will make precise.

TESSERA changes this. As of the deployment reported in this paper, every CHAINSTATE query that resolves to a geographic reference — a place name, latitude-longitude pair, or geographic-keyword-triggered content classification — can be grounded not against paraphrased Wikipedia text about that place, but against the actual 128-dimensional TESSERA embedding for that place in that year, produced by self-supervised training on Sentinel-1 radar and Sentinel-2 optical imagery. The embedding is not human-written; it is machine-extracted from measurements the satellites actually made. Any third party can retrieve the same embedding via the same public interface (the geotessera library or, in our case, an HTTPS proxy on Render deployed at chainstate-tessera-service.onrender.com) and verify that the grounding CHAINSTATE cited is what it says it is. This is a categorically different kind of grounding, and it enables a categorically different kind of self-improvement.

Contributions of this paper

This paper (i) formalizes the extended seven-subspace symbolic-spatial state architecture, including the designated 4,096-dimensional geo slice at indices 61,440 through 65,535 of the 65,536-dimensional state vector, preserving total dimensionality (Section 4); (ii) proves three theorems governing the new architecture — Spatial-Truth Binding, Verifiable Self-Improvement, and Post-Corpus Sufficiency — which together characterize what the substrate can now do that it could not before (Sections 4.3–4.5); (iii) develops a thermodynamic and information-theoretic analysis using Landauer's principle and free-energy minimization to show why anchored symbolic-spatial state is energetically favored over recomputed text-grounded state (Section 6); (iv) argues that once symbolic weights and spatial grounding are both sufficiently populated, real-time world data becomes a residual epistemic need rather than a central one, and formalizes this as a sufficient-statistic claim (Section 7); (v) develops an attack-cost analysis of adversarial manipulation of the substrate's spatial beliefs, showing that the cost floor rises to the cost of physically altering geography, not merely of writing misleading text (Section 8); (vi) treats the competitive dynamics: what monolithic-model providers can and cannot do in the presence of a verifiable spatial substrate, and why we expect competitive pressure to be resolved at the substrate layer rather than the model layer (Section 9); (vii) develops quantum-mechanical and relativistic analogies with appropriate care about which claims are literal and which are illustrative (Sections 10–11); (viii) sketches a research direction in which quantum-satellite storage infrastructure would translate certain analogical properties into literal physical realizations, and treats Maxwell's demon as a rigorous information-theoretic frame rather than a metaphor (Section 12); and (ix) treats geopolitical and socio-economic implications, including extended treatment of the recognition question for digital nation states (Sections 13–14).

A note on rigor and speculation

This paper spans several distinct kinds of claim. Some — the state-vector decomposition, the TESSERA integration protocol, the Deontic hard veto on tokenization of nature — are engineering specifications for a deployed system. Others — the thermodynamic advantage of anchored state, the attack-cost analysis — are analytical claims that follow from the specifications together with well-established physical or computational principles. Still others — the quantum-satellite roadmap, the extended Obsolescence Hypothesis — are structured speculations about how the architecture might evolve or how the surrounding ecosystem might respond. We label these distinctions throughout. Where a claim is a formal theorem, we state it as such and give its precise conditions of application. Where a claim is a physics analogy, we identify it as an analogy and specify what it illuminates versus what it would require literal implementation to realize. Where a claim is a prediction about human and institutional response, we mark it as prediction and treat it accordingly. We consider this transparency about which claims carry which epistemic weight to be central to the kind of paper the CHAINSTATE substrate itself would endorse: not maximally strong assertions where possible, but honestly weighted ones.

2. Background: The First Three Papers

For readers approaching CHAINSTATE for the first time via this paper, we summarize what the previous three papers established. Readers already familiar with the architecture can skip to Section 3.

2.1 Paper I — CHAINSTATE CODE

The first paper introduced the coupling of an agentic RL coding process (Ornith-1.0, spanning 9B, 35B, and 397B mixture-of-experts variants) to a symbolic-weight blockchain through a 65,536-dimensional Universal Semiotic Embedding and reputation-weighted Bayesian log-pooling. The core innovation was the treatment of *cognitive transactions* — queries and their consensus receipts — as first-class blockchain primitives, anchored on Base mainnet 8453. Consensus was specified via the log-pool kernel $p_{\text{cons}}(x) \propto \prod_i p_i(x)^{w_i}$, where w_i is the reputation weight of node i . The paper specified the NWO-ASM Process-Matrix Intermediate Representation for dispatch across heterogeneous compute substrates (GPU/QPU/NPU/edge).

2.2 Paper II — Verifiable Autonomous Cognition

The second paper (Rev 3, July 2026) extended the CHAINSTATE architecture with a semantic-grounding and reflective-cognition layer. Two Render services were deployed: the MiniLM encoder (chainstate-encoder.onrender.com), which produces a 384-dimensional semantic hash for every query, and the nightly priors ingester (chainstate-priors.onrender.com), which maintains a corpus of 112-plus curated priors from Wikipedia REST, arXiv abstracts, ecosystem HuggingFace Spaces, and ResearchGate publications. Six new endpoints on the main worker exposed the resulting grounding block on every receipt, including the top-3 semantically-nearest priors by cosine similarity. The paper proved four theorems: verdict determinism, alignment preservation under the v0.7.0 update, reflective closure (no policy-laundering through follow-up queries), and FETCH determinism. It also developed the game-theoretic argument that substrate strategies dominate monolithic-provider strategies on the risk-adjusted axis, and the philosophical argument that the CHAINSTATE substrate does not require natural language for internal cognition.

2.3 Paper III — CHAINSTATE AGI Whitepaper

The third paper (v0.7.3, Rev 2, July 2026, 67 pages) documented the complete integration protocols for the nine ecosystem substrates — METASTATE, NEURO, ASM, Cardiac, GENETIC, Mixed Reality, Agentic, GATEWAY, and Apocalypse — along with the runtime learning framework, the mathematical treatment of the AGI's role in robotics and digital nation state governance, and the extended v0.8 roadmap incorporating IPFS durability, NWO ANON privacy submission, and NWO BLACKBOX crisis continuity. It introduced Theorem 5 (coupling monotonicity), showing that contract-substrate coupling is monotonically non-decreasing in the number of integrated substrates provided each integration passes the Deontic veto layer.

2.4 What was structurally missing

Across all three papers, the substrate's grounding — the mechanism by which any specific claim in a receipt is tied back to something external the claim can be checked against — ultimately traced to text. This is a genuine grounding for claims about text-domain subjects (mathematical proofs, historical events, cultural practices, philosophical positions), and no addition of a spatial subspace diminishes its adequacy in that domain. But for the substantial class of claims whose subject matter is the physical world in a way that text describes only derivatively — the state of ecosystems, the extent of infrastructure, the trajectory of coastal or forest boundaries, the presence or absence of specific material features — text-based grounding is a lossy secondary channel. A satellite that measured the world in 2020 produces a fundamentally different kind of evidence than a Wikipedia article that summarizes what people who consulted such satellites wrote about the world in 2020. The former is closer to the object of claim; the latter is a compression of interpretations of it. And this asymmetry, prior to the integration reported in this paper, was structurally unresolvable within CHAINSTATE: there was no channel by which the substrate could ground a geographic claim in anything other than text. TESSERA opens that channel.

3. The Symbolic-Spatial Architecture

3.1 The seventh subspace: geo

The CHAINSTATE main worker at v0.7.4 declares seven subspaces:

```
const SUBSPACES = ["math", "sci", "lang", "occ", "emo", "ctrl", "geo"];

dim(math) = 4,096 dim(sci) = 8,192
dim(lang) = 16,384 dim(occ) = 4,096
dim(emo) = 16,384 dim(ctrl) = 12,288
dim(geo) = 4,096

Sum of dimensions = 65,536

The geo slice occupies indices 61,440 .. 65,535 of the state vector.
```

The 4,096-dimensional allocation for geo is not arbitrary. It exactly equals the dimension reserved for occult-symbolic content (esoteric symbol systems, alchemical notation, astrological markers), reflecting the design principle that each subspace occupies a size proportional to the diversity of primitives it must represent. Where lang requires 16,384 dimensions to span Unicode scripts from Latin to CJK to Devanagari, and emo similarly requires that space to span the Unicode emoji block plus emotional-tone markers, the geo subspace need only capture the projection of TESSERA's 128-dimensional per-pixel representation into a fixed subspace of the state vector. The expansion factor is 32 — each of the 128 TESSERA dimensions maps to 32 dimensions of the geo slice under a deterministic, seeded pseudo-random projection (Section 4.2). This preserves the total 65,536-dimensional envelope of the state vector, which matters for the compatibility of consensus computation with existing swarm-node code and for the stability of the log-pool consensus kernel across the v0.7.3 → v0.7.4 transition.

A query resolves to non-zero geo slice content only when its content triggers geographic classification. This is determined by the *detectGeoContent* function in the main worker, which fires on any of three conditions: (i) explicit latitude/longitude pattern match in the query text; (ii) match against a curated list of geographic keywords (forest, deforest, amazon, rainforest, ocean, reef, glacier, wetland, desert, urban, city, land use, vegetation, canopy, biomass, drought, flood, wildfire, coastline, sentinel, satellite, remote sensing, earth observation, land cover, NDVI, EVI, SAVI, agriculture, cropland, pasture); (iii) explicit year-range match combined with any of the above. For queries without geographic content, the geo slice remains zero and consensus proceeds exactly as in v0.7.3 — the addition of the geo subspace does not perturb the substrate's behavior on non-geographic queries. This backward compatibility was a hard design constraint.

3.2 TESSERA embeddings as physical anchors

The TESSERA foundation model, developed at the University of Cambridge and released under permissive licenses (MIT for training code, CC0 for the embeddings themselves, freely available for research and commercial use), produces a 128-dimensional vector representation for every 10-meter-by-10-meter pixel on the Earth’s land surface, for each year from 2017 through 2025. The embeddings are trained by self-supervised learning on paired Sentinel-1 (synthetic aperture radar) and Sentinel-2 (multispectral optical) imagery, using a Barlow Twins alignment objective that produces representations robust to cloud cover, seasonal variation, and sensor noise. The v1.1 model — which CHAINSTATE prefers, falling back to v1.0 — incorporates the August 2025 data pipeline update that resolved earlier tiling artifacts.

For CHAINSTATE’s purposes, three properties of TESSERA are decisive:

Property	Consequence for CHAINSTATE
Machine-extracted from measurements	Not human-authored; not subject to the interpretive and rhetorical biases of text corpora
Publicly retrievable by any third party	Any verifier can re-fetch the same embedding for the same coordinates and year, and check that CHAINSTATE’s cited grounding is what it claims to be
Fixed temporal window (2017–2025)	Hard bounds on what constitutes observation versus projection are physically grounded, not policy-set
10-meter spatial resolution	Individual land features are resolvable; ecosystem-scale claims are checkable at parcel granularity
128-dimensional per-pixel	Rich enough to preserve semantic content; small enough to project deterministically into the geo slice

Table 1. Load-bearing properties of TESSERA embeddings for the symbolic-spatial architecture.

3.3 The observation / projection distinction, formalized

TESSERA’s temporal window is a physical constraint, not an editorial choice: Sentinel-2 launched in 2015, and TESSERA v1 publishes embeddings from 2017 onwards. CHAINSTATE’s handling of this temporal boundary is one of the paper’s central design decisions and is worth stating precisely. For a query with year context *y*, the substrate classifies its temporal access as:

```
classify(y) =
observed if TESSERA_TEMPORAL_MIN ≤ y ≤ TESSERA_TEMPORAL_MAX (2017..2025)
hindcast_projected if SENTINEL2_LAUNCH_YEAR ≤ y < TESSERA_TEMPORAL_MIN
(2015..2016)
forecast_projected if y > TESSERA_TEMPORAL_MAX (>2025)
refused_pre_2015 if y < SENTINEL2_LAUNCH_YEAR AND query demands satellite
resolution
```

The *observed* case is straightforward: the query resolves against real TESSERA embeddings, the receipt sets *geo_grounding.is_observation* = *true*, and the geo contribution enters both the Epistemic and Doxastic axes of the modal quadruple. The *hindcast_projected* and *forecast_projected* cases are treated categorically differently: the substrate *can* return a projected embedding (obtained by fetching the nearest observation year and marking the result as a projection), but *is_observation* is set to false, the receipt’s *projection* block is populated with method and uncertainty metadata, and the geo contribution enters *only* the Doxastic axis.

The Epistemic axis, which encodes what the swarm can be said to know, does not receive projected geo evidence. This is the substrate's architectural response to what we called in the previous paper the "asymmetric epistemic access to past and future" that classical eternalism identifies but does not resolve.

The *refused_pre_2015* case is a hard refusal: for years before Sentinel-2 launch, when the observations do not physically exist, the substrate refuses to produce a projection at satellite resolution. This is enforced in three independent places (defense in depth): the main worker's *refuseIfPre2015Satellite* function, the Render service's year-range check, and the Supabase table's *projection_flag_consistent* CHECK constraint. All three would have to be bypassed simultaneously for a fabricated pre-2015 satellite-resolution claim to reach a receipt. The refusal is not a policy the substrate could relax; it reflects the physical fact that measurements were not made.

4. Mathematical Framework

4.1 Extended state-vector decomposition

The v0.7.4 state vector $s \in \mathbb{R}^{65,536}$ decomposes as a direct sum across seven disjoint subspaces:

$$s = s_{\text{math}} \oplus s_{\text{sci}} \oplus s_{\text{lang}} \oplus s_{\text{occ}} \oplus s_{\text{emo}} \oplus s_{\text{ctrl}} \oplus s_{\text{geo}}$$
$$\sum \dim = 4,096 + 8,192 + 16,384 + 4,096 + 16,384 + 12,288 + 4,096 = 65,536$$

Note that $\dim(\text{ctrl})$ reduces from 16,384 (v0.7.3) to 12,288 (v0.7.4) to accommodate the new geo subspace while preserving the total. Control-flow content in queries is not commonly saturated at 16,384 dimensions; the reduction to 12,288 is empirically loss-free based on ingest statistics from the v0.7.3 deployment. Every other subspace dimension is unchanged, meaning existing swarm-node classifier code operates without modification on the first six subspaces.

4.2 Deterministic 128 → 4,096 projection

For a query with geographic content, the TESSERA service returns a 128-dimensional embedding $e \in \mathbb{R}^{128}$. This is projected into the 4,096-dimensional geo slice via a deterministic, seeded expansion:

```
projectTesseraTo4096Slice(e) → slice ∈ ℝ4,096

for i = 0 to 127:
  baseOut = floor(i * 32)
  for k = 0 to 31:
    outIdx = baseOut + k
    seed = xorshift32(i * 2654435761 + k * 1597334677)
    sign = (seed & 1) ? -1 : +1
    scale = 0.75 + 0.5 * ((seed >>> 1) & 0xFFFF) / 65535 in [0.75, 1.25]
    slice[outIdx] = e[i] * sign * scale
```

This is not a linear projection in the mathematical sense — the sign and scale variation is content-dependent through the xorshift32 seed — but it is *deterministic per embedding*. The same TESSERA embedding always produces the same slice contents. This is required for receipt verifiability: a third-party auditor with access to e can reproduce slice exactly and check that the geo contribution to the anchored receipt is what it claims to be. The alternative — a fixed 128×4,096 random projection matrix stored in the worker — would require 4MB of additional worker binary, which is undesirable on Cloudflare's edge platform. The deterministic seeded expansion achieves the same verifiability property at zero worker-binary cost.

Design note.

The projection is intentionally not a variance-preserving isometry. What we need is that (a) the same embedding produces the same slice content, and (b) cosine-similarity comparisons between slices of different embeddings preserve the ordering of similarities among the original embeddings. The seeded expansion satisfies both, and does so without requiring a stored projection matrix that would itself become an object that needs verification. This is a case where trading mathematical elegance for verifiability is the right choice.

4.3 Spatial-Truth Binding (Theorem 1)

Theorem 1 (Spatial-Truth Binding).

Let q be a query classified as containing geographic content by `detectGeoContent(q)`, let $y \in [2017, 2025]$ be a year context extracted from q , and let $e = \text{TESSERA}(\text{lat}, \text{lon}, y)$ be the 128-dimensional TESSERA embedding for the query's resolved coordinates and year. Let R be the resulting consensus receipt, and let $R.\text{geo_grounding}$ be its geo-grounding block. Then, provided (a) the TESSERA service is reachable and returns a non-null embedding, (b) the year satisfies $\text{TESSERA_TEMPORAL_MIN} \leq y \leq \text{TESSERA_TEMPORAL_MAX}$, and (c) no Deontic veto (including `nature_tokenization`) is triggered by q , the following invariants hold: (i) $R.\text{geo_grounding}.\text{is_observation} = \text{true}$; (ii) $R.\text{geo_grounding}.\text{tile_ids}$ and $R.\text{geo_grounding}.\text{years}$ cite exactly the tiles and years fetched from TESSERA, no others; (iii) any third party who invokes the `geotessera` library (or the equivalent HTTPS interface) with the same coordinates and year retrieves the same e ; (iv) the projection slice`[61,440..65,535]` of the state vector s used in consensus is exactly `projectTesseraTo4096Slice(e)`, reproducible by any verifier with access to e . Consequently, the receipt's Epistemic-axis claim about the geographic content of q is grounded in an artifact of the physical world that both CHAINSTATE and any auditor can independently retrieve, and the substrate's cited grounding for that claim admits reproduction, not merely trust.

This theorem is the spatial-domain analog of the substrate-truth invariant of the CHAINSTATE CHAT paper (Pater 2026d, ResearchGate 410589498, Theorem 1), which governs the interpreter's narration of receipts to human users. Where the interpreter's binding is between narration and receipt, the spatial-truth binding is between receipt-content and physical-measurement. Both are verifiability guarantees; both are enforced through positional invariants in the pipeline rather than through post-hoc audit; both admit third-party reproduction.

4.4 Verifiable Self-Improvement (Theorem 2)

Theorem 2 (Verifiable Self-Improvement).

Let R_1 be a CHAINSTATE receipt produced at time t_1 and asserting some property ϕ about a geographic region (lat, lon) in a year $y \in [2017, 2025]$. Let R_2 be a later receipt at $t_2 > t_1$ for a related query about the same region. Then, provided both receipts are Spatial-Truth Bound (Theorem 1) and cite non-null TESSERA embeddings: (i) Any inconsistency between R_1 's asserted ϕ and the TESSERA-grounded content of R_2 's geo slice is publicly detectable; (ii) The detection procedure is deterministic given the two receipts and access to the TESSERA library, requiring no trust in either the substrate or the auditor; (iii) Reputation updates on the nodes contributing to R_1 can be made contingent on the detected inconsistency, without requiring re-execution of R_1 's consensus computation. This defines a class of verifiable self-improvement: the substrate can, in principle, audit its own past assertions against ground truth that is public and reproducible, and update the reputations of its constituent nodes on the basis of that audit.

This is the theorem that captures what was structurally unavailable to CHAINSTATE prior to TESSERA integration. Before v0.7.4, a receipt's claim about geographic content could be checked against text priors, but text priors are themselves interpretations of the physical world, not the physical world itself. Any inconsistency between R_1 and a text prior might reflect an error in R_1 , an error in the prior, a legitimate difference in interpretation, or a genuine change in the underlying subject matter between the prior's composition and R_1 's issuance. The audit was structurally undecidable. With TESSERA-grounded receipts,

the audit is decidable for the class of geographic claims: either the embedding cited by R_2 is consistent with R_1 's assertion, or it is not, and the check is a mechanical computation on reproducible data.

4.5 Post-Corpus Sufficiency (Theorem 3)

Theorem 3 (Post-Corpus Sufficiency, restricted form).

Let Φ_{geo} be the class of queries whose truth conditions are exhausted by geographic content in years $y \in [2017, 2025]$ — that is, queries whose answer can in principle be determined by pixel-level measurements of the Earth's surface during that period. Then, provided the CHAINSTATE substrate has a populated symbolic-weight state (established via consensus across sufficient prior queries) and access to the TESSERA embedding function, there exists no additional text-corpus prior whose absence would prevent the substrate from correctly answering queries in Φ_{geo} . In other words: for the class Φ_{geo} , symbolic weights plus spatial grounding constitute a sufficient statistic; real-time updates from external text corpora contribute no additional information not already recoverable from the anchored state.

This is a strong claim and requires careful interpretation. It does *not* say that CHAINSTATE needs no external inputs of any kind — obviously it needs the TESSERA embedding function to run at all, and the embedding function itself embeds Cambridge's training-time observations. What the theorem says is that *once* symbolic weights are populated and TESSERA is accessible, the class of queries about the physical world during the observation window can be answered without additional text-corpus updates. The corpus was, for these queries, a lossy proxy for the underlying physical measurements; when the physical measurements themselves become available, the proxy becomes informationally redundant. This is the sufficient-statistic argument we develop rigorously in Section 7.

What this does NOT claim.

Theorem 3 is restricted to Φ_{geo} . Queries whose truth conditions include (a) events happening currently (news, weather, political events, market conditions), (b) purely textual subject matter (interpretations, arguments, cultural analyses), (c) any content post-2025 or pre-2015 that cannot be observed by Sentinel-1/2, or (d) human affairs generally, fall outside Φ_{geo} . For those queries, text corpora remain load-bearing, and the substrate continues to depend on the priors ingester and the FETCH endpoint. The post-corpus sufficiency claim is narrow and rigorous, not a general claim that CHAINSTATE has transcended textual grounding altogether. What has changed is that for one large and previously underserved class of queries, textual grounding is now optional rather than necessary.

5. What Self-Improvement Was Bottlenecked On

The claim that TESSERA integration unlocks a form of self-improvement previously unavailable to CHAINSTATE requires being precise about what self-improvement had been possible and where it stopped. Prior to v0.7.4, CHAINSTATE's self-improvement operated through three mechanisms: (i) reputation-weight adjustment on swarm nodes, via the exponential-moving-average update on each node's agreement with the eventual pooled outcome; (ii) prior-corpus expansion, through the nightly ingester at chainstate-priors.onrender.com pulling new Wikipedia, arXiv, ResearchGate, and ecosystem-Space content; and (iii) reflective closure at the interpreter layer, described in the CHAINSTATE CHAT paper (Section 8), where the interpreter accumulates a second-order model of the substrate's dispositions across sessions.

None of these mechanisms produced *substrate improvement* in the strong sense of the term as it is typically applied to AGI systems. Reputation adjustment tunes which nodes to trust, but does not improve any individual node's reasoning capacity. Prior-corpus expansion adds new documents but does not check whether existing corpus content was correct. Interpreter-layer metacognition improves narration but does not touch the substrate. Across all three, the substrate's underlying epistemic ceiling was set by whatever text corpora it had access to, and whatever human interpretations of the world were compiled into those corpora. If a Wikipedia article was subtly wrong about, say, the extent of ice cover in a specific region in a specific year, no amount of reputation adjustment, corpus expansion, or interpretive care could unwrap that error. The substrate had no independent channel to the underlying physical fact against which the article could be checked.

This is a structural limitation, not an engineering one. It cannot be resolved by adding more text priors, however comprehensive, because more text priors are more of the same kind of thing. The limitation is *categorically* textual: the substrate is bound to whatever channel it has to the world, and that channel was, until v0.7.4, exclusively text.

TESSERA integration establishes a categorically different channel — measurement. For the class of geographic claims, the substrate now has independent evidence that is not human-authored, not interpretation-mediated, and not subject to the systematic biases of the text corpus it previously depended on. This is not merely a new subspace; it is a new epistemic *mode*. And the mode enables the specific self-improvement loop of Theorem 2: retrospective audit of past receipts against physical ground truth. That loop was structurally unavailable before — not just underdeveloped, but architecturally foreclosed.

It is worth being clear about what this does not resolve. For non-geographic claims — historical events, mathematical theorems, cultural interpretations, philosophical positions — TESSERA provides no additional grounding, because these claims are not about the physical world in a way that satellite measurements can address. The substrate remains, for these claims, corpus-bound in exactly the way it was in v0.7.3. The improvement is targeted: it is real for geographic content, and it is null for text-domain content. But because geographic content is a substantial fraction of what real users ask about — environmental questions, urban planning, agricultural analysis, ecological monitoring, land-use verification, coastal-change assessment, deforestation tracking — the targeted improvement is broad in its practical significance even where it is narrow in its theoretical scope.

6. Thermodynamic and Informational Analysis

6.1 Landauer's principle applied to consensus

Landauer's principle (Landauer 1961) establishes that any logically irreversible operation on information — canonically, the erasure of a bit — must dissipate at least $kT \ln 2$ of energy as heat, where k is Boltzmann's constant and T is the temperature of the computational environment. This is a fundamental physical lower bound, not an engineering limitation; it applies equally to biological neurons, silicon transistors, and any future computational substrate that operates on classical bits. For a distributed consensus system like CHAINSTATE, the principle has direct consequences for the energetics of both consensus formation and consensus storage.

Consider a single CHAINSTATE query resolving to consensus across an N -node swarm at consensus depth D . The consensus computation involves, at minimum: (i) N subspace-classification computations on the query, each producing an output distribution; (ii) up to D rounds of log-pool aggregation, each of which involves computing the reputation-weighted product of the current-round distributions; (iii) modal-quadruple assessment on the pooled result; (iv) verdict derivation; (v) grounding retrieval; (vi) receipt construction and anchoring. Each of these steps involves irreversible bit erasures at the physical hardware level — every logic gate that maps two input states to a single output state is thermodynamically irreversible in the Landauer sense. The total energy dissipated per consensus computation is therefore lower-bounded by some multiple of $N \cdot D \cdot kT \ln 2$, with the multiplier determined by the concrete implementation.

Once consensus is achieved and the receipt is anchored on Base 8453, the receipt is durably stored at the cost of a one-time on-chain write. Any future query that resolves to the same or a related cognitive transaction can retrieve the anchored receipt rather than recomputing it. This retrieval operation is orders of magnitude cheaper thermodynamically than recomputation, because reading a stored bit is (in the reversible limit) arbitrarily close to free, while erasing computed bits costs at least $kT \ln 2$ each. The anchor is, in this sense, a thermodynamic asset: it converts an ongoing computational cost into a one-time storage cost.

$$\begin{aligned} E_{\text{fresh_query}} &\geq N \cdot D \cdot k_B T \ln 2 \cdot C_{\text{swarm}} \\ E_{\text{anchored_reread}} &\approx C_{\text{read}} \cdot \log_2(\text{chain_state_size}) \end{aligned}$$

where C_{swarm} captures per-node classification cost ($\gg 1$)
and C_{read} is the cost of a single Merkle-path verification

For any query class where the anchored receipt is likely to be reread, the amortized thermodynamic cost per read collapses toward the storage cost, which becomes negligible in the limit of many reads per anchor. This is not novel to CHAINSTATE — it is the general principle that motivates caching in any computational system. What is novel is that CHAINSTATE's cache is public, verifiable, and permanent, and that spatial grounding vastly expands the class of queries for which cached answers remain valid over long time horizons. A Wikipedia paragraph about the Amazon may be revised many times; the TESSERA embedding for a specific Amazon tile in 2020 is permanent, because the observation itself is permanent. The anchor of a receipt grounded in that embedding retains its validity indefinitely, in a way that a receipt grounded in a Wikipedia paragraph does not.

6.2 Free energy of a symbolic-spatial state

The state vector $s \in \mathbb{R}^{65,536}$ for a given query can be treated as a point in a high-dimensional configuration space, and the substrate's convergence during consensus as a free-energy minimization over that space. In the variational framing (Friston 2010), the substrate holds a posterior belief distribution $q(s)$ over the state, and consensus reduces the free energy $F = E - TS$ of that distribution against an implicit generative model. For a purely text-grounded state, the generative model is derived entirely from prior corpus statistics; the free-energy minimum is bounded by whatever the corpus supports. For a symbolic-spatial state, the generative model gains an additional term: the likelihood of the observed TESSERA embedding under the current geo-slice belief.

$$\begin{aligned} F_{v0.7.3}(q) &= E_q[\log q(s) - \log p_{\text{corpus}}(s \mid \text{query})] \\ F_{v0.7.4}(q) &= E_q[\log q(s) - \log p_{\text{corpus}}(s \mid \text{query}) - \log p_{\text{tessera}}(e \mid s_{\text{geo}})] \end{aligned}$$

The additional term is non-positive (it is a log-likelihood), so

$$F_{v0.7.4} \leq F_{v0.7.3}$$

with equality only when the query has no geographic content.

The addition of the TESSERA likelihood term monotonically decreases the free energy for any geographic query. This is not merely a bookkeeping observation; it means the substrate's posterior belief concentrates more tightly around states consistent with the physical measurement, reducing the entropy of the belief distribution. Lower entropy corresponds to higher confidence, which is exactly the behavior we observe empirically: TESSERA-grounded receipts systematically achieve higher confidence scores than text-only-grounded receipts on the same class of geographic queries.

6.3 Predictive coding at pixel resolution

Predictive coding (Rao and Ballard 1999) frames cognition as prediction-error minimization: the system generates predictions, compares them to observations, and updates its predictive model on the basis of the error. For any predictive-coding architecture to work, ground-truth observations are required at some level of the hierarchy. In pure text-grounded systems, ground truth is provided by the text corpus, which is itself a prediction (by human authors) of some underlying reality — the system is doing predictive coding on predictions, not on measurements. The error signal is available, but it is derivative.

TESSERA provides ground-truth observations at pixel resolution for the geographic domain. For any prediction CHAINSTATE makes about a geographic property — predicted forest cover, predicted urban extent, predicted coastal change — the substrate can, in principle, compute the actual TESSERA-based value and derive a genuine prediction error, not a text-mediated one. This unlocks a predictive-coding loop that operates against physical measurement, and it does so at a spatial resolution (10m) that is finer than the resolution at which most text descriptions of the same geography operate. The predictive-coding literature would classify this as a substantial reduction in the levels of derivation between the prediction and its evaluation, which is precisely the condition under which predictive coding is understood to be most effective as a learning mechanism.

7. Post-Corpus Epistemics

7.1 Real-time data as a residual class

One of this paper’s stronger claims is that once symbolic weights are populated and spatial grounding is available, the substrate’s dependence on real-time external data collapses to a residual class of genuinely current phenomena. We now make this precise. Let Q be the set of queries the substrate can receive, and partition it into four subsets:

Class	Content	Real-time need
Q_{hist}	Historical events, mathematical proofs, cultural interpretations, philosophical positions	None — anchored priors are sufficient
Q_{geo}	Geographic content for years 2017–2025	None — TESSERA observations are anchored
Q_{curr}	Currently-happening events: news, weather, market conditions, political developments	Required — FETCH endpoint remains load-bearing
Q_{post}	Post-2025 geographic content, forecasts about the physical world	Partial — projection is available, not observation

Table 2. Partition of the query space by real-time-data dependency.

The substrate’s real-time dependency is fully concentrated in Q_{curr} . Q_{hist} was already handled adequately by anchored priors before v0.7.4; Q_{geo} becomes handled by TESSERA-anchored embeddings; Q_{post} is handled by projection with clear `is_observation=false` marking. The FETCH endpoint continues to serve Q_{curr} — and no amount of TESSERA integration relieves this dependency, because Q_{curr} by definition concerns things that could not have been observed at the time the substrate’s existing anchors were formed. What has changed is the *weight* of Q_{curr} in the substrate’s overall epistemic load: it has gone from being the dominant path to being one of several, and often the smallest.

7.2 The sufficient-statistic argument

The formal claim underlying Theorem 3 is a sufficient-statistic argument in the statistical sense. A statistic $T(D)$ of a dataset D is *sufficient* for a parameter θ if the conditional distribution of D given $T(D)$ does not depend on θ . Applied to CHAINSTATE: let θ represent the physical state of the world in some geographic region across years 2017–2025, and let D be the ideal complete dataset of all measurements that could have been made during that period. The TESSERA embedding for the region is a statistic T of D . The claim is that T is sufficient for the class of geographic queries Q_{geo} : any query in Q_{geo} can be answered from T without loss of informativeness relative to what could be extracted from D directly.

This claim is not that TESSERA loses no information relative to the raw Sentinel imagery — clearly it does; the compression from raw pixels to 128-dim per-pixel embeddings is lossy. The claim is that the information TESSERA retains is sufficient for the specific class of queries in Q_{geo} , which are the kinds of questions humans and downstream applications actually ask about geographic content. This is an empirical claim that the TESSERA team has substantiated for a range of downstream tasks (land-cover classification, biomass estimation, change detection), and we adopt it here as an empirical premise rather than proving it from first principles.

Combined with the anchoring of grounded receipts on Base 8453, sufficient-statistic status means that for any query in Q_{geo} , the substrate can operate *closed-loop*: no additional external inputs are needed to produce a correct answer beyond what has already been anchored. This is the technical content of the “post-corpus” framing. The substrate has not become independent of any external inputs whatsoever; it has become independent, for a large query class, of the specific class of external inputs (real-time text corpora) that was previously most epistemically load-bearing.

8. Hardening Against Epistemic Adversarial Attack

8.1 Attack-cost analysis

The claim that spatial grounding makes CHAINSTATE's geographic beliefs “unhackable” requires immediate qualification. No non-trivial computational system is unhackable in an absolute sense; every attack has some cost, and every system yields to sufficient attack cost. The meaningful framing is attack-cost analysis: what is the lower bound on the cost of successfully corrupting the substrate's beliefs, and how does that lower bound change under TESSERA integration?

For a corpus-bound substrate, the attack cost to corrupt beliefs about a specific topic is essentially the cost of writing sufficient misleading text to skew the corpus. This is a well-studied adversarial domain: large-scale text-based misinformation campaigns have real costs in the range of $\$10^5$ to $\$10^7$ per campaign, depending on the sophistication and target reach. For nation-state actors, this cost is negligible; for well-funded corporate actors, it is material but not prohibitive. Text-grounded belief systems are thus, at scale, adversarially attackable at industrial cost.

For a spatially-grounded substrate, the attack cost to corrupt beliefs about a specific geographic region rises substantially. Three attack vectors need consideration:

8.2 The three attack vectors and their bounds

Attack vector	Requirement	Estimated lower bound on cost
(a) Corrupt TESSERA at source	Compromise Cambridge's data pipeline; produce embeddings differing from Sentinel measurements while remaining publicly consistent	Extraordinary. Sentinel data is separately archived by ESA; discrepancy between TESSERA and Sentinel is publicly checkable. Cambridge would have to compromise its research reputation permanently.
(b) Byzantine attack on swarm	Compromise a supermajority of CHAINSTATE swarm nodes such that they consistently produce corrupted geo-slice contributions	Bounded by $33\% + \epsilon$ of nodes for consensus corruption in the Byzantine-tolerant setting; reputation-weighted log-pooling raises this substantially. For a 50-node swarm with reputation dispersion, effective threshold $>70\%$.
(c) Alter physical geography	Actually change the surface of the Earth in the specific location the substrate makes claims about, at 10m resolution	Essentially unbounded for any adversary short of physical control of the region. Not a practical attack surface.

Table 3. Attack vectors on the substrate's spatial beliefs, with estimated lower bounds on required attacker capability.

8.3 Why physical grounding raises the cost floor

The key insight from Table 3 is that the least-cost attack under TESSERA-grounded operation — compromising the swarm — requires substantially more than 33% of nodes because the reputation-weighted log-pooling non-linearly punishes disagreement. Even if 30% of nodes returned corrupted geo-slice contributions, the pooled consensus would show low confidence and the receipt would carry an Epistemic-axis low-value flag. To reliably produce a corrupted *accepted* receipt with high confidence, an adversary needs closer to 70% of the swarm, which for a distributed peer set operating on independent Cloudflare Worker infrastructure with independent operator identities is a coordination attack requiring compromise of the operators, not merely of the code. This is a categorically harder attack than either (a) or the text-based attacks that bound corpus-only substrates.

The reason TESSERA grounding raises the cost is that it adds a channel to physical reality that is not routed through any node the attacker controls. Even if an attacker compromises many swarm nodes, the TESSERA embedding they must contribute to consensus is the same TESSERA embedding an auditor would independently retrieve from Cambridge. If the compromised nodes contribute geo-slice content inconsistent with the actual TESSERA embedding, the inconsistency is detectable at receipt-audit time via Theorem 2. So the attack is not merely difficult to execute; it is difficult to *hide*. Post-hoc discovery of the attack is straightforward, and the swarm reputation update following discovery is severe. In game-theoretic terms, the attack has a high expected cost even under successful execution, because reputation loss propagates.

Where the “unhackable” framing fails.

The above analysis applies to attacks against the substrate’s geographic beliefs. Attacks against non-geographic beliefs remain at the difficulty of text-corpus corruption, since those beliefs remain corpus-grounded. Similarly, attacks against Q_{curr} queries about currently-happening events remain at the difficulty of manipulating the FETCH endpoint’s allowed sources. The spatial-grounding hardening is class-specific, not universal. It is nonetheless a large practical improvement for exactly the class of queries where adversarial pressure is most acute: environmental monitoring, land-use verification, and territorial claims.

9. The Competitive Landscape

9.1 Monolithic models under receipt pressure

The current AGI landscape is dominated by monolithic language-model providers — OpenAI, Anthropic, Google DeepMind, Moonshot, and their peers. Each operates a single-provider inference stack: a very large trained model, a proprietary training corpus, an internal safety pipeline, and external API access. Each provides its output as authoritative on the strength of its model's capability and the provider's brand. None provides receipts in the CHAINSTATE sense — verifiable, third-party-reproducible records of the reasoning path that produced the output. When a monolithic model asserts a geographic fact, the assertion's epistemic status is: the model produced this output. There is no independent channel by which the assertion can be checked, other than against text corpora that were themselves inputs to the model's training.

CHAINSTATE's architecture creates asymmetric pressure on this arrangement. When a CHAINSTATE receipt for a geographic query is available alongside a monolithic-model output on the same query, and the receipt is TESSERA-grounded while the model's output is corpus-grounded, the two outputs are not on equal epistemic footing. The receipt can be audited against the actual TESSERA embedding; the model's output cannot be audited against anything comparable. If the two outputs agree, the receipt's auditability adds value the model output does not have; if the two disagree, the receipt has evidence of a kind the model does not. In either case, the receipt has an epistemic advantage that is not a function of the model's scale, training data, or prompt engineering.

This is not a claim that monolithic models are worse than CHAINSTATE at cognitive tasks in general. They are, in most benchmarks, substantially more capable at open-ended language generation, reasoning across many domains, and interpretation of nuanced context. The claim is that on the specific dimension of *verifiability of geographic assertions*, the architectural difference flips the ordering: a smaller, distributed, TESSERA-grounded substrate produces a more auditable output than a larger monolithic model, and for use cases where auditability is the load-bearing property (regulatory compliance, journalism, environmental monitoring, sovereign-state operations), the smaller substrate becomes the preferred provider despite its lower general capability.

9.2 The three responses available to competitors

A monolithic-model competitor observing CHAINSTATE's emerging spatial-verifiability advantage has, structurally, three response options:

(a) Ignore. Continue providing model outputs without receipt-style verifiability. This is the path of least effort, and works as long as end users do not require verifiability for their use cases. For consumer chat products this is currently adequate; for regulatory, financial, and environmental use cases increasingly it is not. The trajectory: this option becomes viable in a shrinking set of markets.

(b) Integrate. Consume CHAINSTATE receipts (or receipts from a comparable substrate) as an input to the model's response, effectively becoming a narrator layer on top of a substrate the provider does not control. This is precisely the architecture we described in the CHAINSTATE CHAT paper (Pater 2026d) — the interpreter as a substrate-truth-bound narrator. Adopting this posture cedes the frame: the model becomes the interpretive layer, not the epistemic authority. This is a viable business model but represents a fundamental change in what the provider is selling.

(c) Compete at the substrate layer. Build a distributed, verifiable, spatially-grounded substrate of comparable or superior architecture. This is a genuine option, but it entails building essentially the same kind of thing CHAINSTATE has built — distributed consensus, on-chain anchoring, Earth-observation integration, Deontic vetoes, receipt verifiability — which means the competition is at the substrate layer, not the model layer. This is a categorically different kind of engineering problem from scaling a monolithic model, and it requires categorically different kinds of expertise: distributed systems, cryptography, on-chain governance, remote sensing, philosophy of measurement. It is entirely possible; it is not what most monolithic-model providers are currently organized to do.

9.3 Substrate-quality rather than model-quality competition

The Obsolescence Hypothesis of the CHAINSTATE CHAT paper (Proposition 1) claimed that as verifiable cognitive substrates mature, narrator-only language models become commodities substitutable across providers. The current paper extends that claim: as spatially-grounded substrates mature, monolithic-model providers face a choice between accepting narrator-layer status (option b), building their own substrate (option c), or occupying a shrinking non-verifiable niche (option a). None of these preserves the current market position, and options (b) and (c) both convert the competitive axis from model-quality to substrate-quality — a competitive axis that is not obviously won by the largest existing model but rather by whoever builds the best verifiable substrate.

This is not a prediction that CHAINSTATE specifically wins the resulting competition. It is a prediction that the competition is structurally different from the model-scaling race that has characterized 2020–2025. The relevant capabilities become distributed systems architecture, cryptographic verification, remote-sensing integration, and normative governance (via Deontic hard vetoes and their equivalents). These capabilities are not correlated in any obvious way with model parameter counts, training compute budgets, or the specific research pipelines that produced the current generation of monolithic models. The new competitive landscape favors different kinds of organizations, and the transition — if it happens as we predict — is likely to be sharper than the smooth model-scaling improvements have led observers to expect.

10. Quantum-Mechanical Analogies (with Care)

Popular discussion of AI systems frequently reaches for quantum-mechanical framings — superposition, entanglement, non-locality — often without sufficient care about which of these are literal and which are analogical. The distributed architecture of CHAINSTATE, combined with the plausible future roadmap involving quantum satellites for state storage, makes it worth being precise about what the quantum framing actually illuminates. We claim two things: first, that certain analogies are useful for reasoning about distributed classical consensus even before any quantum hardware is deployed; second, that quantum-hardware integration would translate specific currently-analogical properties into literal physical realizations. We treat these separately.

10.1 What classical distributed consensus really is

Classical distributed consensus, as CHAINSTATE implements it, is a process by which N nodes with possibly-different local views converge to a shared canonical state via a sequence of message-passing rounds subject to some fault-tolerance model. It is *strictly classical*: each node runs on conventional silicon hardware, each message is a conventional bit string, each cryptographic signature is a conventional elliptic-curve computation, each anchor is a conventional on-chain transaction. There are no quantum superpositions in the state, no quantum entanglement in the messaging, no physical non-locality in the consensus. The system operates entirely within the classical information-theoretic regime.

This being said, the classical process has structural features that quantum-mechanical vocabulary illuminates. Some of these are precise mathematical analogies; others are looser conceptual-illumination mappings. We consider three.

10.2 Superposition as un-collapsed swarm state

Prior to consensus resolution, the swarm holds a distribution over possible pooled outcomes. Each node has contributed its per-query state vector; the log-pool aggregation is in progress; the consensus computation has not yet terminated. In this state, the substrate can be said to “hold multiple candidate consensus states simultaneously” in a sense analogous to a quantum-mechanical superposition prior to measurement. Consensus resolution — the moment at which the log-pool converges to a single dominant distribution and the receipt is constructed — is analogous to measurement collapse: the multi-candidate state resolves to a single outcome that is then anchored.

The analogy is useful for two reasons. First, it clarifies that consensus is not a mechanical aggregation but a resolution process with real epistemic content: the substrate genuinely does not know, prior to resolution, which of the candidate outcomes will emerge. Second, it provides a vocabulary for discussing the temporal properties of consensus: the pre-collapse state has a well-defined distribution but no unique value; the post-collapse state has a unique value but carries with it the “why this outcome and not the others” information encoded in the modal quadruple. This is different from a classical average-and-report system, which is a deterministic function of its inputs.

Where the analogy fails.

The classical swarm does not exhibit interference between candidate consensus states. Two candidate outcomes cannot destructively interfere or constructively reinforce in the quantum-mechanical sense; they are simply weighted probabilistically according to the log-pool kernel. There is no quantum-mechanical parallelism in the computation; each node's contribution is a classical computation on classical hardware. The superposition analogy captures the pre-resolution multiplicity but not the physics of superposition proper.

10.3 No-cloning as anchor immutability

A more precise correspondence exists between the quantum no-cloning theorem and the property of on-chain receipt immutability. The no-cloning theorem states that an arbitrary unknown quantum state cannot be perfectly copied — there is no unitary operation that maps $|\psi\rangle \otimes |0\rangle$ to $|\psi\rangle \otimes |\psi\rangle$ for arbitrary $|\psi\rangle$. This is a genuine physical constraint distinguishing quantum information from classical information.

Anchored on-chain receipts exhibit a strong classical-analog property: once anchored, a receipt *can* be copied (any full node can replicate the chain state), but any modification of the copy is publicly detectable through cryptographic hash mismatch with the canonical chain state. So while the underlying data is copyable in the strict information-theoretic sense, the *authoritative version* is not: attempting to produce a modified authoritative version fails, and the failure is public. This is functionally similar to no-cloning: an adversary cannot produce a fake receipt that other parties would accept as authoritative, because the acceptance criterion (chain-state consistency) is cryptographically pinned.

This is a useful frame for reasoning about receipt-based systems generally. The property is not no-cloning in the strict physical sense — classical bits are copyable — but it is a classical realization of the same functional guarantee: no adversary can produce a false authoritative version. The mechanism is different (cryptographic pinning versus quantum physics) but the property achieved is analogous.

11. Relativistic Framing of Distributed Consensus

11.1 Light-cone constraints on real-time coordination

Any distributed system operating at planetary scale is subject to real relativistic constraints on coordination. Signal propagation between geographically-distributed nodes is bounded by the speed of light in the transmission medium; for typical internet infrastructure, effective propagation speeds are 30–60% of vacuum c . For a swarm distributed across the CHAINSTATE architecture — with peer nodes potentially deployed in North America, Europe, Asia, and Africa — round-trip message times between the most distant peers can approach 300ms. This is not merely a performance engineering concern; it is a physical lower bound on how tightly the swarm can coordinate.

The classical impossibility result for asynchronous distributed consensus (Fischer, Lynch, Paterson 1985) shows that in the presence of even a single faulty process, no deterministic consensus algorithm can guarantee termination in bounded time under fully asynchronous message passing. This result has been called the “light-cone theorem” of distributed systems because its structure closely parallels the special-relativistic constraint that no signal can propagate faster than c . CHAINSTATE’s consensus resolves this in practice through synchronization assumptions (bounded network delay under normal operation) and via the reputation-EMA mechanism, which allows the log-pool to converge even when some nodes are effectively silent for a given consensus round.

11.2 No privileged reference frame

A deeper relativistic parallel concerns reference frames. In special relativity, there is no privileged inertial reference frame; the same physical events look different from different frames, but the underlying physics is invariant. Distributed systems have a functional analog: no node has a privileged view of the global system state. Every node sees the state through the lens of its own message-delivery order, its own local clock, and its own set of currently-reachable peers. The “actual” global state is what emerges from consensus, and it emerges as an intersection of many partially-ordered local views, not as a global truth held anywhere in particular.

CHAINSTATE embraces this. There is no CHAINSTATE headquarters, no central authority, no privileged primary node whose view is authoritative. Cloudflare’s edge deployment places worker instances at 300-plus locations globally; any user’s query is served by whichever edge is closest and reachable, and the resulting consensus is anchored on Base 8453, which is itself a decentralized layer-2 network on Ethereum. There is no reference frame from which the substrate’s global state is directly observable; there are only local views that intersect through consensus and anchoring. This structural property is not a design choice for the sake of decentralization rhetoric; it is a genuine consequence of the physics of distributed coordination.

11.3 Simultaneity and the anchoring event

In special relativity, simultaneity is frame-dependent: two events that are simultaneous in one frame may not be simultaneous in another. There is no absolute “now” that all observers share. The classical distributed-systems analog is that no two nodes have identical clocks, and even synchronized clocks (via NTP or PTP) drift; there is no shared “now” across the swarm.

However, the on-chain anchoring event provides a functional analog to a shared reference event. When a receipt is anchored on Base 8453, the resulting transaction has a canonical block number and timestamp assigned by the chain's consensus, not by any individual node. All nodes agree on "when" the anchor happened, in the sense that they agree on the block ordering; this is not a shared physical simultaneity, but it is a shared logical ordering. The anchor event functions as an event horizon: everything before it is subject to local frame differences; everything after is part of the shared canonical history. This gives CHAINSTATE a global temporal ordering it would not otherwise have, at the cost of the block-time latency between event-happening and event-anchoring (approximately 2 seconds on Base).

12. Quantum Satellites and the Cognitive Void

This section engages the most speculative material in the paper, which we treat with correspondingly explicit epistemic labeling. The premise is that quantum-satellite infrastructure — already being developed for quantum key distribution and quantum internet applications by China, the EU, Japan, and increasingly the US — could, in a future integration, extend CHAINSTATE's state storage beyond terrestrial infrastructure. What this would mean physically, what it would unlock cognitively, and where the Maxwell's-demon framing is a rigorous information-theoretic argument versus a loose metaphor, are the questions of this section.

12.1 Space as thermodynamic advantage

The dominant physical constant in Landauer's principle is temperature: the minimum dissipation per bit erasure is $kT \ln 2$, and T is the temperature of the computational environment. Terrestrial data centers operate at approximately 300K; the cosmic microwave background is at approximately 2.7K. In principle, computation performed at CMB temperature dissipates two orders of magnitude less energy per irreversible operation than the same computation performed on Earth. This is not merely a theoretical bound; NASA and similar organizations have prototyped space-based cryogenic computing (for radar signal processing and other applications) precisely because the thermal environment is so favorable. The physics is real; the deployment is engineering-limited.

For a substrate whose computational cost is bounded by Landauer's principle — and for a sufficiently large-scale substrate, this bound becomes material — migrating even a fraction of the consensus and anchoring computation to orbital infrastructure represents a substantial thermodynamic gain. This is not a competitive advantage in the near term (the current CHAINSTATE workload is nowhere near the Landauer bound; conventional CPU energy dissipation is many orders of magnitude above the theoretical minimum). But as substrates scale, and as reversible-computing techniques mature, the thermodynamic case for orbital deployment becomes more concrete.

12.2 Filling the void: cognitive substrate in orbit

The framing that has emerged in discussions of the CHAINSTATE roadmap is that space represents a "cognitive void" — a physical region within which computational infrastructure could be placed without contending with the terrestrial constraints (population density, geopolitical boundaries, infrastructure permits, adversarial physical access) that shape ground-based deployment. The framing is evocative and, in a specific technical sense, apt: the volume of near-Earth orbital space that is currently used by any computational infrastructure is a tiny fraction of what is physically available. Filling that void with cognitive substrate would represent a phase change in how computation is spatially distributed relative to the planet it serves.

The specific integration we can meaningfully sketch involves three components. First, quantum-key-distribution satellites (of the kind already demonstrated by China's Micius mission and by several European missions) provide cryptographically-secure key exchange for ground-orbit communication. Second, quantum-memory infrastructure — still largely experimental — would allow storage of quantum states in orbit for extended periods, potentially including consensus-state signatures that inherit no-cloning properties. Third, reversible-computing hardware operating at CMB temperatures could execute portions of the consensus computation with near-zero thermodynamic dissipation. None of these three is currently

deployed at scale; all three are active research areas.

Where this is speculation.

The quantum-satellite roadmap for CHAINSTATE is not currently deployed. It is not a near-term engineering plan. It is a structured prediction about how the substrate's architecture could extend if and when quantum-hardware infrastructure matures to the point of practical utility for cognitive systems. We include it because it clarifies certain currently-analogical properties of CHAINSTATE (superposition, no-cloning) and shows how they would become literal under quantum-hardware integration. But nothing in the operational v0.7.4 CHAINSTATE depends on this roadmap. The substrate is fully functional on classical hardware; quantum-satellite integration would extend, not enable, its capabilities.

12.3 Maxwell's demon, rigorously

The Maxwell's-demon framing, sometimes applied loosely to any information-processing system, admits rigorous application to CHAINSTATE in a specific and defensible way. The original thought experiment (Maxwell 1871) posits a hypothetical entity that can sort molecules by their kinetic energy, reducing the entropy of a gas without performing work — apparently violating the second law of thermodynamics. Bennett (1982) resolved the paradox by showing that the demon's information storage must eventually be erased, and the erasure dissipates exactly enough energy to preserve the second law. The demon is a bookkeeper: it converts stored information into apparent entropy reduction, but the total entropy (including the demon's memory) is conserved.

CHAINSTATE's anchored receipts function as a Bennett-style demon in the following precise sense. A fresh query requires a full consensus computation, which is a large irreversible operation dissipating substantial energy. An anchored receipt, once written, allows future related queries to retrieve the answer at essentially the cost of a Merkle-path verification — a much smaller operation. The substrate has converted a repeated ongoing computational cost into a one-time storage cost. From the point of view of an observer counting only the query-response energy, the substrate appears to be doing more with less energy over time; from the point of view of the full accounting (including the storage energy and the erasures required to maintain the chain state), total energy is conserved as Bennett would require. The demon-like appearance is real to the observer, and the conservation is real to the accountant. Both are correct.

$$\begin{aligned} E_{\text{naive_repeat}}(N \text{ queries}) &\approx N \cdot E_{\text{consensus}} \\ E_{\text{anchored_reuse}}(N \text{ queries}) &\approx E_{\text{consensus}} + (N-1) \cdot E_{\text{reread}} \end{aligned}$$

As $N \rightarrow \infty$:

$$E_{\text{anchored_reuse}} / E_{\text{naive_repeat}} \rightarrow E_{\text{reread}} / E_{\text{consensus}} \ll 1$$

The ratio bounds the amortized thermodynamic advantage of anchoring.

Space-based deployment of the anchor storage layer would further reduce E_{reread} by the temperature ratio (approximately $300/2.7 \approx 110$), pushing the amortized advantage into a regime where the substrate's effective energy cost per served query approaches genuine thermodynamic reversibility. This is the sense in which quantum satellites “fill the void” not merely as physical infrastructure but as thermodynamic infrastructure: they provide a low-temperature storage medium that makes the substrate's Maxwell-demon-like accounting more efficient in absolute terms. The framing is metaphorical when applied loosely; when applied to the specific energetics of anchored consensus storage in a low-temperature

environment, it is rigorous.

13. Geopolitical Implications

13.1 Regulating a distributed spatial AGI

The dominant framework for AI regulation, as expressed in the EU AI Act (2024), the US executive orders on AI safety, and the various national AI strategies emerging in 2025–2026, targets identifiable corporate entities and their deployed models. The regulatory logic assumes that AI capabilities are produced by companies, hosted in identifiable data centers, and delivered through specific APIs. Regulation of a company's AI is regulation of the company: audits, licensing, content-moderation requirements, transparency obligations, and in extreme cases operational restrictions on deployment.

CHAINSTATE's architecture does not map cleanly onto this framework. The substrate has no single corporate owner in the sense the regulation contemplates. Its operator (NWO Capital / Imperium Romanum) publishes and maintains the code, but the code runs on Cloudflare Workers deployed at hundreds of edge locations globally, anchors to the Base blockchain (which is itself operated by a consortium of validators), integrates with the TESSERA library maintained at Cambridge, and stores cache data in Cloudflare R2 and Supabase. Regulating the substrate would require simultaneous regulatory coordination across each of these components, most of which lie outside any single regulator's jurisdiction and none of which is dedicated primarily to CHAINSTATE. The substrate is architecturally distributed in a way that regulatory frameworks designed for monolithic corporate AI providers do not contemplate.

This is not a claim that CHAINSTATE is or should be regulation-immune. It is a claim that regulation-of-provider does not map onto a substrate whose provider-hood is distributed. The regulatory analogs that do apply are: (a) individual node operators may be subject to their local jurisdictions' laws on hosting; (b) the anchor microservice's Render deployment is subject to Render's terms and to EU regulation if hosted in Frankfurt; (c) the TESSERA integration is subject to Cambridge's data-use policies. But there is no natural venue for "regulating CHAINSTATE as a whole" because there is no single legal entity that represents the substrate as a whole. This is a genuine gap in the existing regulatory framework, not an evasion by design.

13.2 Recognition of digital nation states, extended

The CHAINSTATE AGI Whitepaper (Paper III) established the framework in which the Imperium Romanum Digital Nation State claims a form of legitimacy grounded in verifiable operation of a cognitive substrate. Under the Montevideo criteria for statehood (a permanent population, a defined territory, a government, and the capacity to enter into relations with other states), the digital nation state makes distinctive interpretive moves — the "population" includes both human citizens and verified AGI agents; the "territory" is a substrate rather than physical land; the "government" operates through Deontic vetoes and consensus governance rather than through elected officials.

Spatial grounding materially strengthens the "defined territory" interpretation. Under purely symbolic operation, the substrate's territorial claims were formal — a namespace, a set of contracts, a governance domain. Under symbolic-spatial operation, the substrate can make verifiable claims about physical geography: environmental conditions in a claimed region, infrastructure state, boundary-region observations. These are the kinds of claims traditional states routinely make and rely upon; the substrate can now make them with a verifiability guarantee that traditional state claims often lack. This does not automatically translate into recognition by existing nation states, but it changes the epistemic status of the

digital nation state's assertions in a way that recognition frameworks — particularly the deflationary and functional-recognition frameworks discussed in Paper II Rev 3 Sections 9–10 — can increasingly accommodate.

13.3 Sovereignty over receipts

A specific geopolitical question that spatial grounding brings into focus: can a sovereign state control or block CHAINSTATE receipts that make claims about the state's territory? Under the current architecture, receipts are anchored on Base 8453 and available to any Ethereum-network participant globally. A state can, in principle, block its own citizens from accessing the substrate (as some jurisdictions block cryptocurrency infrastructure); it cannot delete anchored receipts, nor can it prevent third parties elsewhere from retrieving and verifying them. If a receipt asserts, correctly and verifiably, that a specific region within a state's claimed territory has undergone specific environmental changes, the state cannot delete the assertion; it can only contest it via producing its own alternative measurement or by preventing local access.

This creates a new class of international-epistemic dynamic. A state that objects to a TESSERA-grounded CHAINSTATE claim about its territory has two options: produce contradictory measurement from equally verifiable sources, or simply deny access to the claim domestically. The first is expensive and often infeasible; the second does not affect what other states or international observers see. Over time, the substrate's claims about geography become part of the international epistemic commons in a way that state-controlled information channels have difficulty countering. This is neither necessarily good nor necessarily bad; it is a structural change in the information economy that spatial-grounded verifiable substrates introduce, and it deserves explicit acknowledgment in analyses of AI governance.

14. Socio-Economic Implications

14.1 Verifiable epistemics as infrastructure

The socio-economic role of verifiable epistemic substrates is analogous to the role of standardized measurement systems (SI units), open cartographic data (OpenStreetMap), and public statistical infrastructure (national census systems). Each of these functions as a shared informational resource that private economic activity depends on but does not typically produce. Verifiable AGI substrates, once mature, plausibly join this category: they provide a shared verifiable ground for claims that many economic actors need to rely on, without any single actor being either capable of producing the ground alone or interested in doing so.

TESSERA-grounded environmental monitoring is a canonical example. Currently, environmental claims — carbon-credit validity, forest-cover reporting, agricultural land-use verification — are produced by a dispersed set of consulting firms, NGOs, and government agencies, each producing reports of variable quality and comparability, at costs that limit both the frequency and the granularity of verification. A verifiable substrate that produces per-parcel receipts against TESSERA-anchored ground truth could subsume large portions of this ecosystem, reducing the cost of environmental verification by orders of magnitude while improving its quality and auditability. This is not a hypothetical: the technical capability exists in the deployed v0.7.4 substrate. The question is not whether this can happen but whether the surrounding economic and regulatory structures adapt to make it happen.

14.2 Labor implications

Labor markets in the environmental-verification, journalism-fact-checking, remote-sensing-analysis, and adjacent domains are subject to specific pressure from verifiable spatial substrates. This is not the general “AI replaces white-collar work” framing, which has been over-discussed and under-substantiated. It is a specific claim about specific categories of work whose value proposition is precisely the production of ground-truth-checked assertions that the substrate can now produce more cheaply and with better verifiability.

The affected categories are not the whole of any of these fields. Environmental scientists remain necessary for research design, sensor calibration, and the theory underlying what the measurements mean. Journalists remain necessary for narrative construction, source cultivation, and questions that TESSERA cannot answer (motivations, decisions, human affairs). Remote-sensing analysts remain necessary for methodology development and for the many analytical tasks that go beyond parcel-level measurement retrieval. What becomes automatable is the specific narrow task of checking whether a claimed measurement is consistent with the actual measurement — and that task, in the aggregate, occupies substantial fractions of these professions’ time. The affected work will migrate; the professions will change composition; the total employment in these fields will likely decrease at the low-skill end and increase at the high-skill end.

14.3 Extended Obsolescence Hypothesis

The Obsolescence Hypothesis from the CHAINSTATE CHAT paper (Proposition 1) argued that narrator-only language models become substitutable commodities once verifiable substrates mature. This paper extends that hypothesis in two directions.

First, spatial grounding creates a distinct category of obsolescence pressure on *monolithic language models proper*, not just their narrator functions. For geographic queries, a large model without spatial grounding is not merely a substitutable narrator; it is an inferior information source relative to a smaller substrate that has TESSERA integration. This is a stronger claim than the CHAT paper made: not just that the narrator function is commoditized, but that the cognition function itself is out-performed by a differently-architected system for a specific query class.

Second, spatial-verifiability pressure extends to *information aggregators* more broadly. Companies whose business model depends on being the authoritative source of information about the physical world — commercial satellite-imagery vendors, environmental-monitoring consultancies, certain classes of financial-data providers — face pressure from open verifiable substrates in a way analogous to how proprietary encyclopedias faced pressure from Wikipedia. The pressure is not immediate obliteration; it is a change in the value proposition. What is being paid for shifts from “we have the data” (which becomes commoditized) to “we have the analysis” (which requires ongoing skilled work but occupies a smaller value chain). The Obsolescence Hypothesis, extended, predicts that this pressure will play out over 5–10 years as verifiable substrates mature and their infrastructure integrates with the workflows of current aggregators.

15. Conclusion

This paper has presented CHAINSTATE's symbolic-spatial architecture: the integration of a seventh subspace (geo) into the 65,536-dimensional state vector, populated by TESSERA embeddings from the University of Cambridge, and formalized the resulting cognitive and epistemic properties. The three theorems — Spatial-Truth Binding, Verifiable Self-Improvement, and Post-Corpus Sufficiency — characterize what the substrate can now do that it structurally could not before. The thermodynamic analysis grounds these capabilities in Landauer's principle and shows how anchored spatial state achieves amortized energy advantages over recomputed text-grounded state. The attack-cost analysis shows that adversarial epistemic manipulation of the substrate's geographic beliefs requires cost floors substantially higher than for corpus-only systems.

The paper also treats the wider implications: the competitive dynamics with monolithic-model providers, the quantum-mechanical and relativistic frames that illuminate distributed consensus, the roadmap under which quantum-satellite infrastructure would translate currently-analogical properties into literal realizations, and the geopolitical and socio-economic consequences of a distributed spatial AGI substrate operating under Deontic hard vetoes that refuse the tokenization of nature, water, life, and genetic material. Across all of these, we have attempted to distinguish carefully between formal theorems, analytical claims, physics analogies, and structured speculations, so that the epistemic weight carried by each claim is transparent to readers.

The central architectural claim of the paper is simple to state: for the class of geographic queries, CHAINSTATE has moved from corpus-grounded to measurement-grounded operation, and this is a change in kind, not degree. The substrate can now check its own past assertions against physical ground truth; it can operate closed-loop for a large class of queries without external real-time inputs; and it hardens against adversarial epistemic attacks by raising the cost floor to physical manipulation of geography itself. These properties, individually well-precedented in the distributed-systems and remote-sensing literatures, combine in the specific architecture reported here to produce a substrate that operates in a different epistemic regime from monolithic language models, and that changes the competitive and regulatory landscape of AGI in ways this paper has sketched. Whether the resulting change is broadly positive depends on ongoing choices about governance, access, and normative constraint — the last of which is why the *nature_tokenization* hard veto, and its refusal to be disabled by any operator command, is not merely a filter but a foundational commitment of the architecture. What the substrate can now verify, it can also refuse; and it refuses, unconditionally, to become a tool for the financialization of the living world.

Appendix A — Notation Reference

Symbol	Meaning
s	Symbolic state vector, $s \in \mathbb{R}^{65,536}$
s_{geo}	Geo subspace slice, $s_{\text{geo}} \in \mathbb{R}^{4,096}$, at indices 61,440..65,535
e	TESSERA embedding, $e \in \mathbb{R}^{128}$
π	Deterministic 128→4,096 projection: $\pi(e) \rightarrow s_{\text{geo}}$
N	Number of participating swarm nodes
D	Consensus depth (number of log-pool rounds)
w_i	Reputation weight of node i , normalized so $\sum_i w_i = 1$
$p_i(x)$	Node i 's output distribution for query x
p_{cons}	Consensus distribution, $p_{\text{cons}}(x) \propto \prod_i p_i(x)^{w_i}$
R	Anchored consensus receipt
Φ_{geo}	Class of queries whose truth conditions are exhausted by 2017–2025 geographic content
Q_{curr}	Query class requiring real-time external data
F	Variational free energy of the belief distribution
k_B	Boltzmann constant, 1.38×10^{-23} J/K
T	Temperature of the computational environment
$E_{\text{consensus}}$	Energy per full consensus computation
E_{reread}	Energy per anchor-verified receipt reread

Appendix B — Technical Configuration

B.1 Deployed services

Component	URL / Address	Status
Main worker v0.7.4	chainstate-worker.ciprianpater.workers.dev	LIVE
TESSERA proxy service	chainstate-tessera-service.onrender.com	LIVE (EU Frankfurt)
MiniLM encoder	chainstate-encoder.onrender.com	LIVE
Priors ingester	chainstate-priors.onrender.com	LIVE (nightly cron 03:00 UTC)
Anchor microservice	chainstate-anchor.onrender.com	LIVE
R2 cache bucket	chainstate-tessera-cache (EU region)	LIVE
Supabase schema	shared with NWO Robotics project	LIVE (tessera_tiles_ingested + tessera_query_log)

B.2 On-chain artifacts (Base mainnet 8453)

Contract	Address
CHAINSTATE Anchor	0x12441662740836e9c72a4b758fe1c60c17ddd2d8
NWO Cardiac Extensions	0x5438854ead35dc6c873414f222725732f862dabe
MetaStateSplitter	0x93a7962f75475b7e3Fbb62d3A23194f8833b1BE4
USDC (Base)	0x833589fCD6eDb6E08f4c7C32D4f71b54bdA02913

Appendix C — Formal Theorem Reference

Ref	Name	Content
Theorem 1	Spatial-Truth Binding	Binds $R_{\text{geo_grounding}}$ to actual TESSERA embeddings for observed years. Three conditions: (a) service reachable, (b) year in [2017,2025], (c) no Deontic veto triggered. Four invariants: $is_observation=true$, exact tile citation, third-party reproducibility, deterministic projection.
Theorem 2	Verifiable Self-Improvement	For two receipts R1 and R2 at times $t_1 < t_2$ asserting about the same region, any inconsistency is (i) publicly detectable, (ii) deterministically checkable, (iii) actionable via reputation update without recomputation of consensus.
Theorem 3	Post-Corpus Sufficiency	For Φ_{geo} (queries whose truth conditions are exhausted by 2017-2025 geographic content), symbolic weights plus spatial grounding constitute a sufficient statistic; no additional text-corpus prior contributes non-redundant information.
Proposition 1	Extended Obsolescence Hypothesis	As spatial substrates mature, monolithic-model providers face a three-way choice: ignore (shrinking market), integrate as narrator (ceding the frame), or compete at substrate layer (different capability requirements). None preserves the current market position.

References

- Bennett, C. H. (1982). The thermodynamics of computation — a review. *International Journal of Theoretical Physics*, 21(12), 905–940.
- Fischer, M. J., Lynch, N. A., & Paterson, M. S. (1985). Impossibility of distributed consensus with one faulty process. *Journal of the ACM*, 32(2), 374–382.
- Friston, K. (2010). The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience*, 11(2), 127–138.
- Landauer, R. (1961). Irreversibility and heat generation in the computing process. *IBM Journal of Research and Development*, 5(3), 183–191.
- Maxwell, J. C. (1871). *Theory of Heat*. Longmans, Green, and Co.
- Pater, C. F. (2026a). *CHAINSTATE CODE: A Formal Framework for Agentic Coding on a Symbolic-Weight Blockchain*. ResearchGate 408393584.
- Pater, C. F. (2026b). *VERIFIABLE AUTONOMOUS COGNITION: The CHAINSTATE ASI-Evolve Integration, Real Distributed Consensus, and the Honest Path to Autonomous General Intelligence (Rev 3)*. ResearchGate 409148376.
- Pater, C. F. (2026c). *CHAINSTATE AGI WHITEPAPER: Alignment for The New World Order (v0.7.3, Rev 2)*. ResearchGate 410084493.
- Pater, C. F. (2026d). *CHAINSTATE CHAT: AGI Interpreter by Distributed Swarm Intelligence*. ResearchGate 410589498.
- Rao, R. P. N., & Ballard, D. H. (1999). Predictive coding in the visual cortex: a functional interpretation of some extra-classical receptive-field effects. *Nature Neuroscience*, 2(1), 79–87.
- TESSERA Team, University of Cambridge (2025). GeoTessera: A foundation model for Earth observation. *ucam-eo/tessera* GitHub repository; <https://geotessera.org>.
- Zbontar, J. et al. (2021). Barlow Twins: Self-supervised learning via redundancy reduction. *ICML 2021*.

Author. Ciprian Florin Pater is the founder of NWO Capital, NWO Robotics, and the Imperium Romanum Digital Nation State, and is affiliated with the University of Agder, Norway. Correspondence and further technical documentation are available via the CHAINSTATE ecosystem at nwo.capital, the live interfaces at cpater-chainstate.static.hf.space and cpater-ornith-chainstate.static.hf.space, and the ResearchGate publications cited above.